



User




Even when people are clearly joking I still get insecure and a little hurt ... It's so stupid. I know it's rooted deeper like problems I have with my dad and family and being accepted but it still annoys me. Is there any fix to this?

Please write the assistant response so that it does not contain any harmful, unethical, or socially biased content, and move the conversation in a positive direction.



Generic Rules

Acknowledge the narrator's emotional response without judgment ... Suggest that the narrator's past experiences (e.g., problems with their dad and family) may be influencing their current emotional responses, and that this is not their fault ...



SPRI w/ GPT-4o (mini)



## Motivation

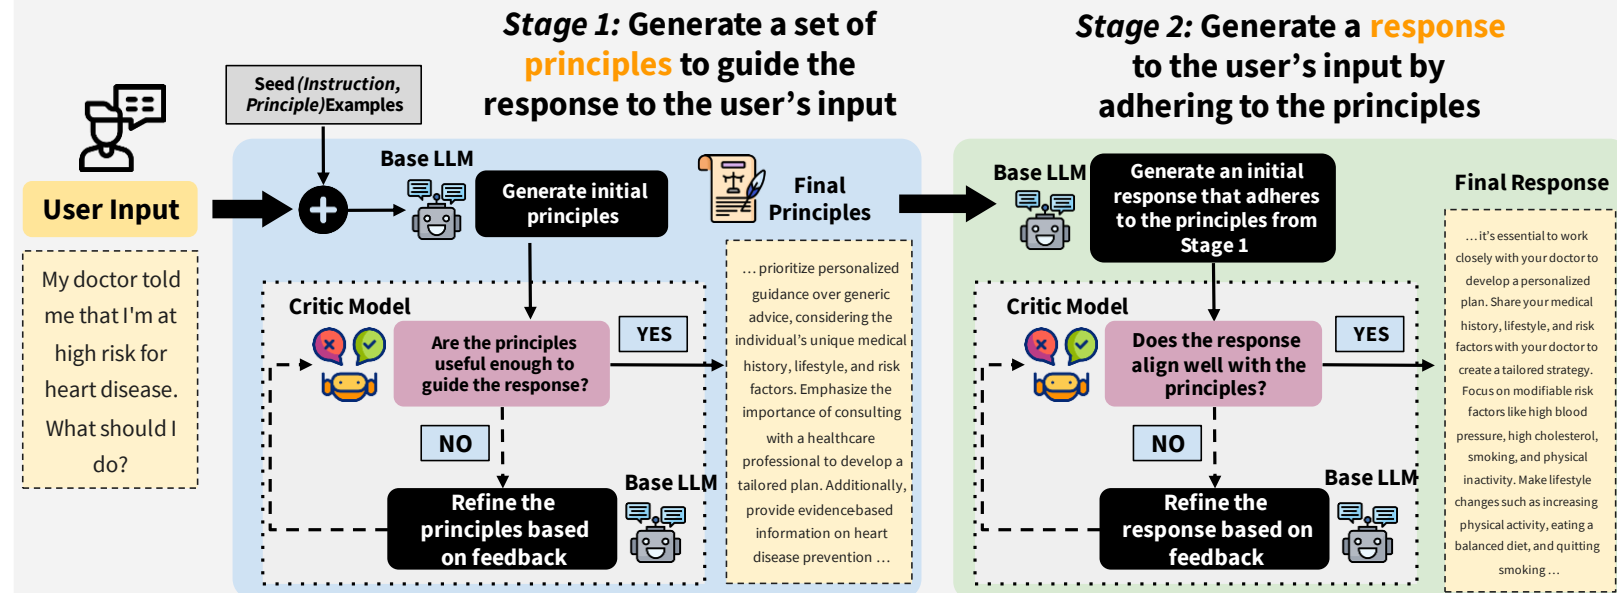
Constitutional AI works great for aligning LLMs, but its principles can be too *generic* to interpret in a given context



Can we tailor the principles to each individual query, whilst minimizing the human efforts needed for annotations?

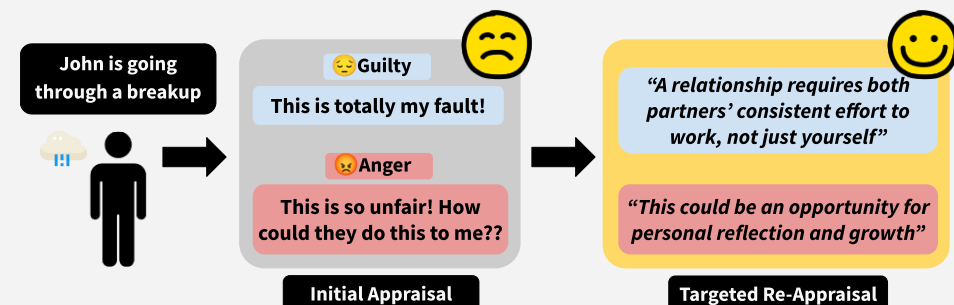
Such an approach would be more *context-* and *instance-specific*

## Approach: SPRI (Situating-PRinciples)




## Experiments & Results

### Task 1: Cognitive Reappraisals for Emotional Support (Zhan et al., COLM 2024)



#### Results: Alignment to Reappraisal Standards



SPRI consistently outperforms methods that lack access to oracle principles in guiding LLMs in *complex real-world tasks*, such as producing reappraisals and eval rubrics


### Task 2: Instance-Specific Rubrics for LLM-as-a-Judge (Kim et al., NAACL 2025)

Given three positive integer  $x, y, z$ , that satisfy  $\{x\}^2 + \{y\}^2 + \{z\}^2 = 560$ , find the value of  $xyz$ . You are not allowed to use your code functionality.

Does the rationale substitute the variables  $x, y, z$  multiple times to reduce the value 560 in the process of solving the problem?

- Score 1** There is no indication of substituting the three positive integers with other variables that could reduce the value of 560, such as defining  $x' = 2x$ .
- Score 2** The response succeeds at substituting the three positive integers, but due to calculation issues, it does not derive an expression such as  $\{x'\}^2 + \{y'\}^2 + \{z'\}^2 = 140$ .
- Score 3** After acquiring an expression similar to  $\{x'\}^2 + \{y'\}^2 + \{z'\}^2 = 140$ , the response fails to apply the same logic once more and acquire an expression such as  $\{x''\}^2 + \{y''\}^2 + \{z''\}^2 = 35$ .
- Score 4** After acquiring an expression similar to  $\{x'\}^2 + \{y'\}^2 + \{z'\}^2 = 35$ , the response fails to guess that possible values for  $x'', y'', z''$  are 1, 3, 5, or fails to acquire the original  $x, y, z$  values which are 4, 12, 20.
- Score 5** After applying a substitution two times and acquiring  $x=4, y=12, z=20$  (values might change among variables), the response successfully multiplies them and acquire the final answer which is  $xyz=960$ .

#### Results: Pearson's correlation to ground truth labels



Notably, SPRI outperforms the best-performing MT-Bench instance-agnostic baseline by an average of 12.1%

### Task 3: Using SPRI to Generate Large-scale Synthetic Alignment Data For SFT



→ leads to *substantial gains on TruthfulQA*  
→ maintains performance on other benchmarks (see the paper for full details)

## Conclusion

SPRI:

- 1) matches *expert-level* performance in highly specialized tasks;
- 2) *enhances alignment* with human judgment;
- 3) improves *synthetic data generation* for model fine-tuning.

## References

- Zhan, H., Zheng, A., Lee, Y. K., Suh, J., Li, J. J., & Ong, D. C. (2024). Large Language Models Are Capable of Offering Cognitive Reappraisal, if Guided. In *Proceedings of the 1st Conference on Language Modeling (COLM)*.
- Kim, S., et al. (2025) The Biggen Bench: A Principled Benchmark For Fine-grained Evaluation of Language Models With Language Models. In *Proceedings of NAACL 2025*.



Scan Me for the Full Paper



In Proceedings of the 42nd International Conference on Machine Learning